Name: Bhavesh Kewalramani

Roll No.: A-25

Section: A

Batch: A1

Practical – 1

Questions:

1. Press the Explorer button on the main panel and load the **weather dataset** and answer the following questions

1. How many instances are there in the dataset?

Ans: 14

2. State the names of the attributes along with their types and values.

Ans:

S.No.	Name of the Attribute	Type of the Attribute
1.	Outlook	Nominal
2.	Temperature	Nominal
3.	Humidity	Nominal
4.	Windy	Nominal
5.	Play	Nominal

3. What is the class attribute?

Ans: Play

4. How will you determine how many instances of each class are present in the data

Ans: By selecting the attribute in the left side panel. In the right side panel we can see
all the information.

5. What happens with the Visualize All button is pressed?

Ans:

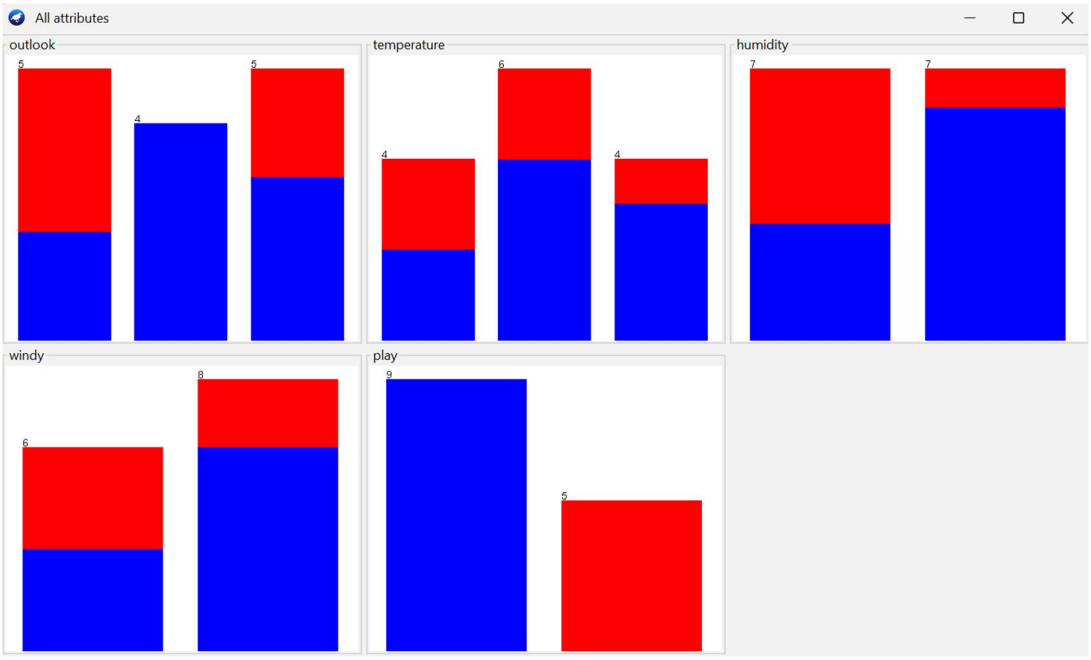

6. How will you view the instances in the dataset? How will you save the changes?

Ans: Using the Edit button in the top

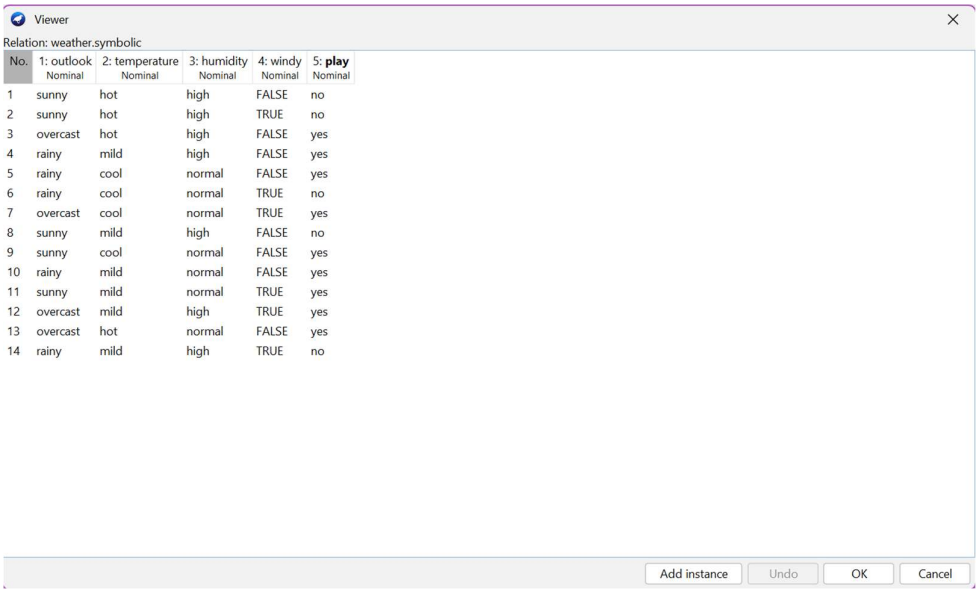

7. Now, extend the dataset to include 50 instances in total.

Ans:

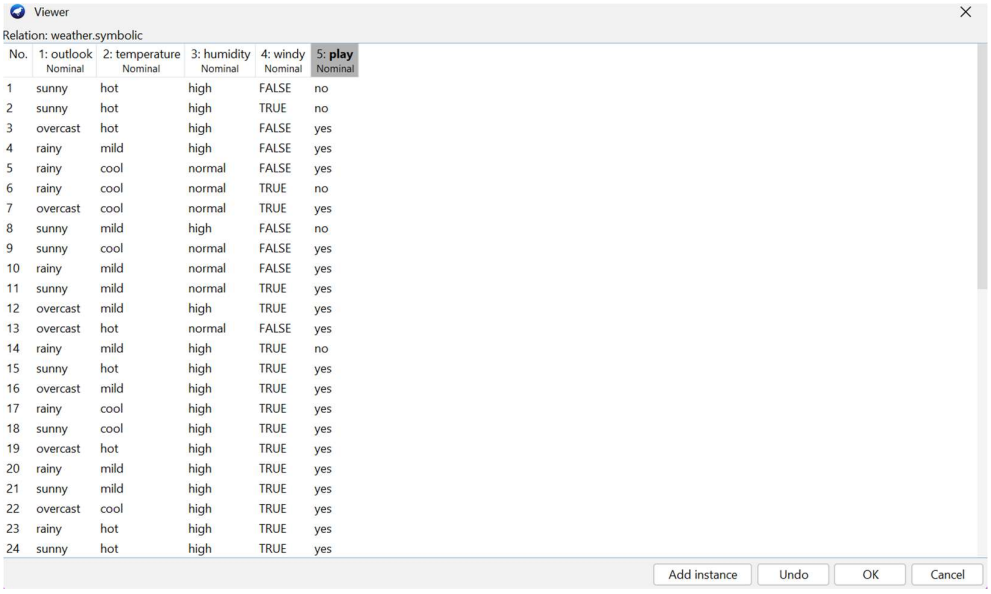

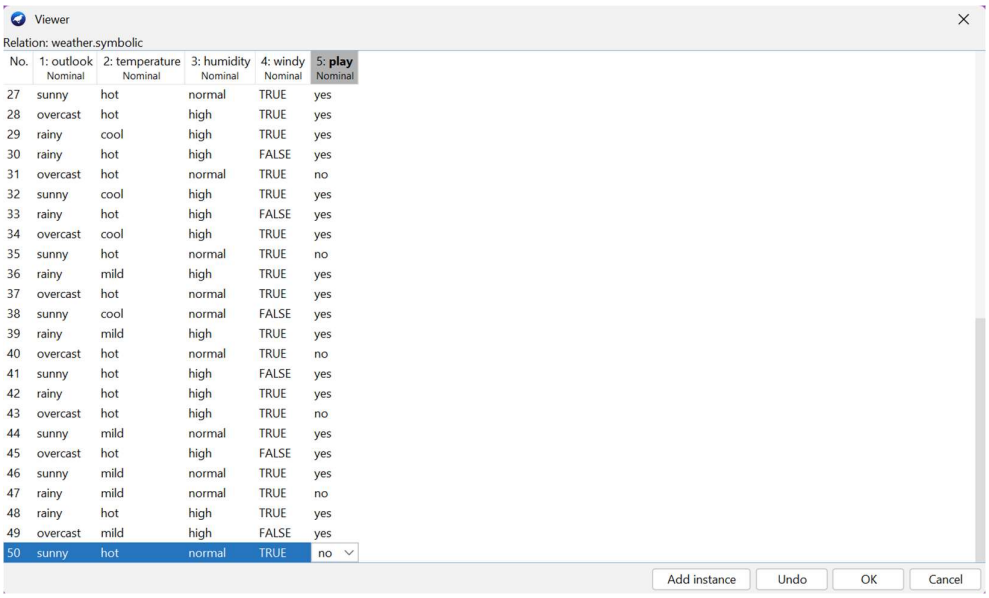

2. Do as directed to apply Filter

1. Use the unsupervised filter RemoveWithValues to remove all instances where the attribute ‘humidity’ has the value ‘high’? Undo the effect of the filter.

Ans:

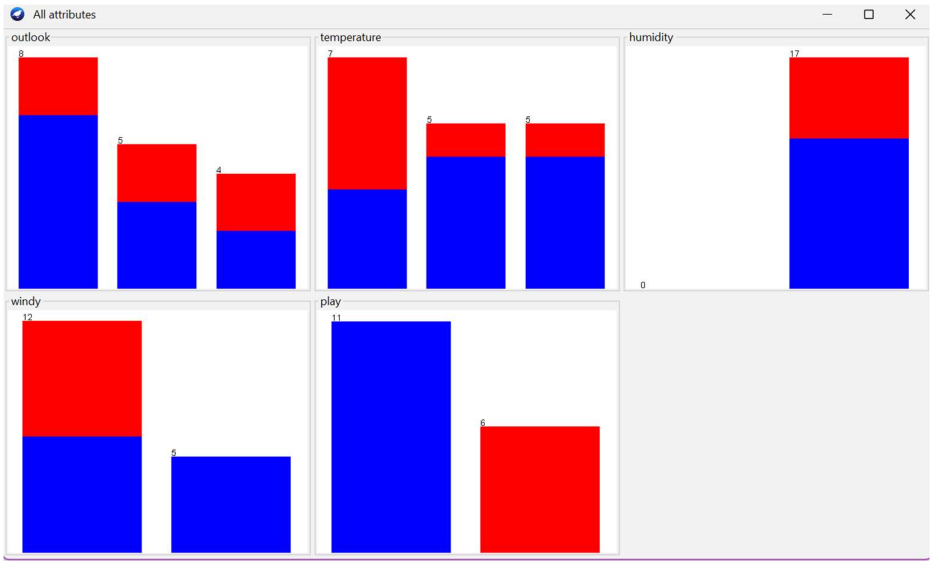

2. Remove the ‘FALSE’ instances of windy attribute and undo the effect.

Ans:

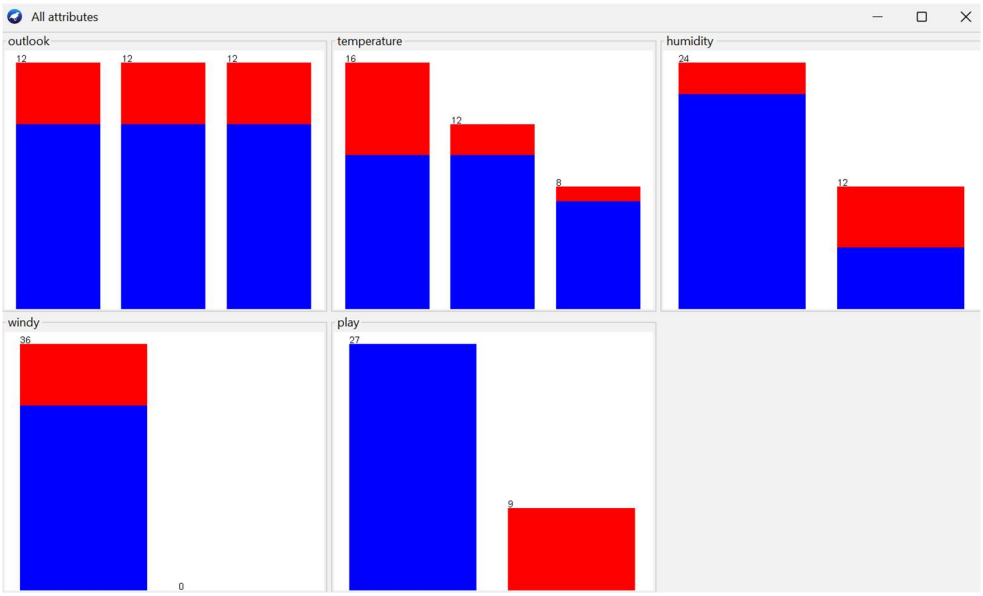

3. Remove the attribute outlook and undo the effect.

Ans:

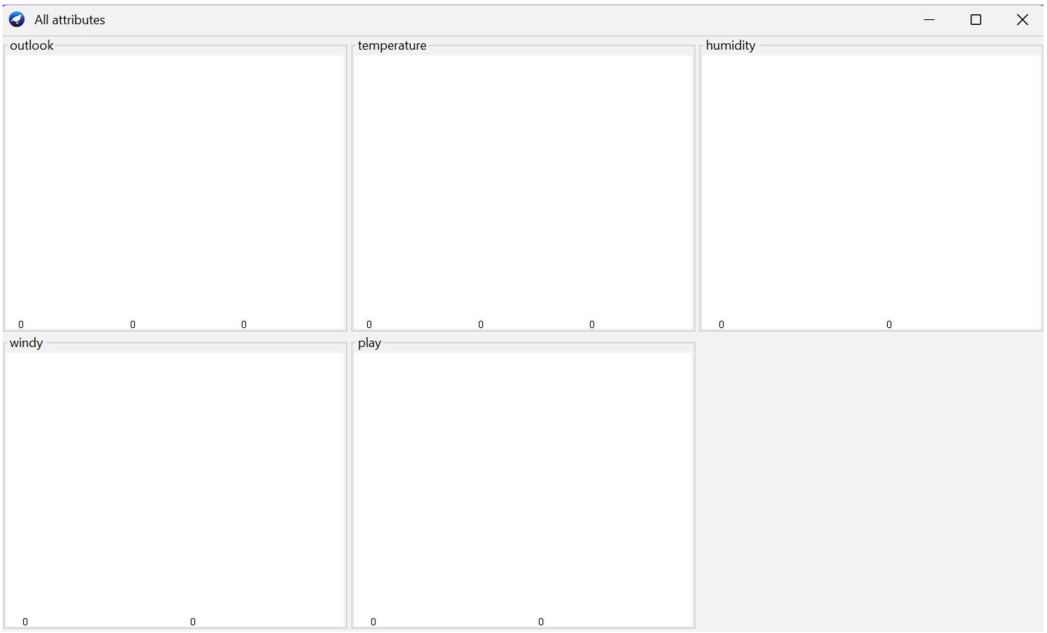

4. Experiment with different filters and report their effects.

Ans:

Resample Filter

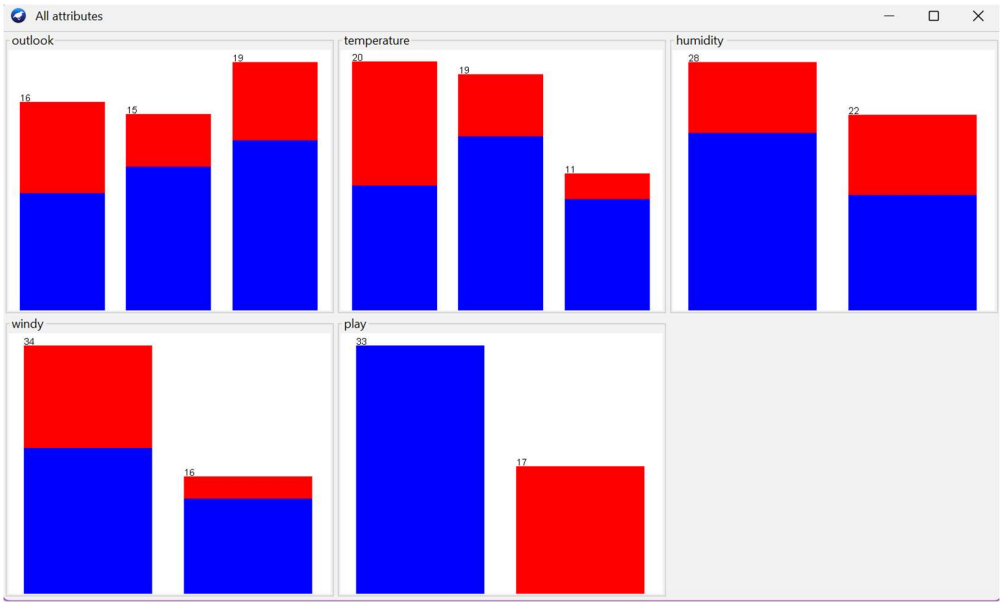

ReservoirSample Filter

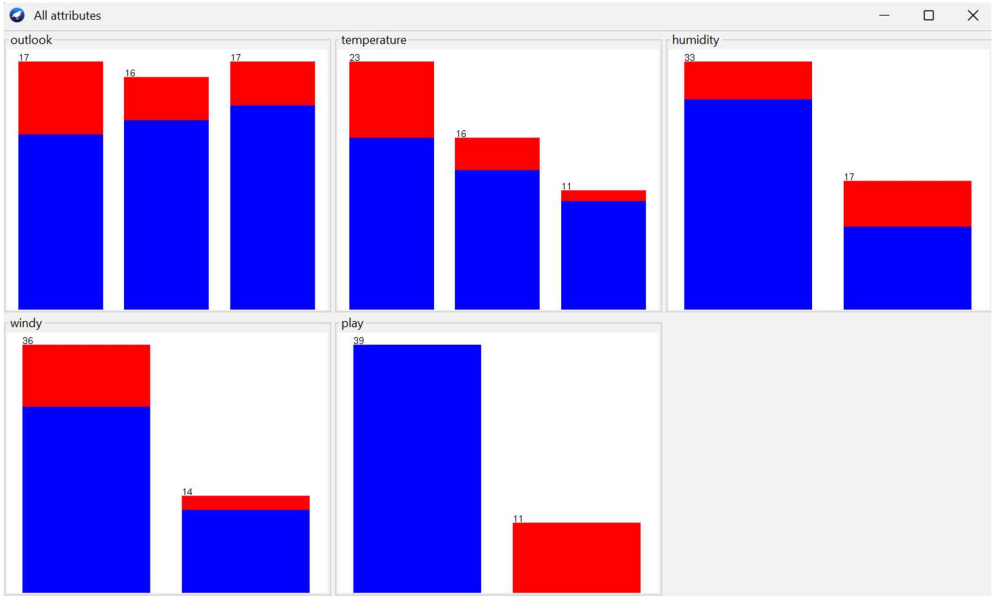

RemoveDuplicates Filter

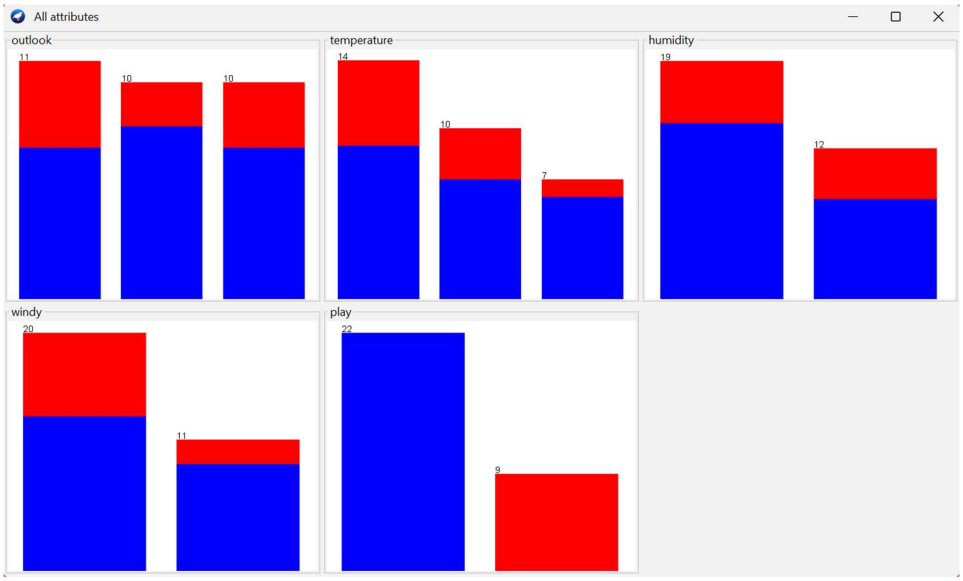

3. Application of Discretization Filters [use sick.arff dataset]

1. Load the 'sick.arff' dataset.

Ans:

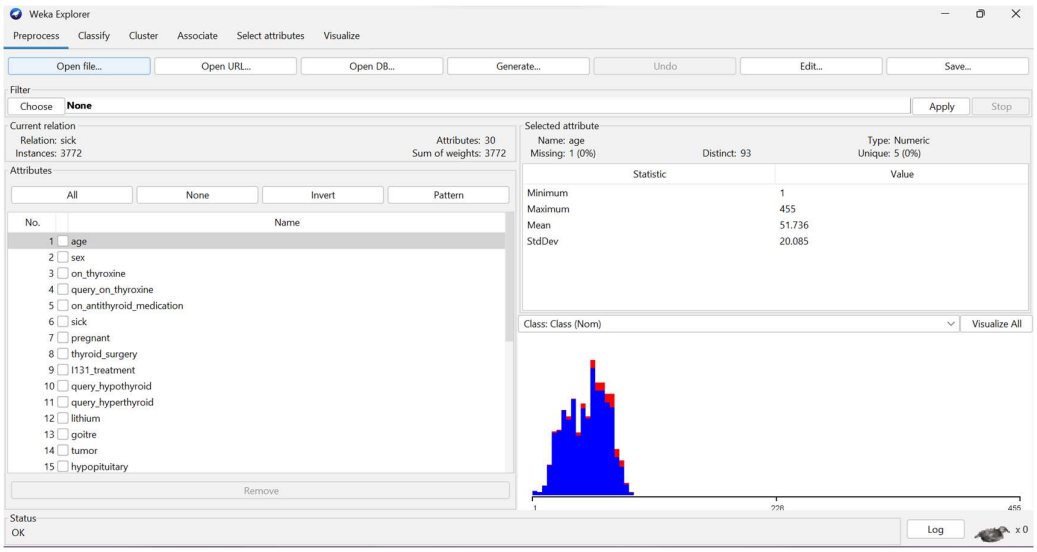

2. Apply the supervised discretization filter on different attributes.

Ans:

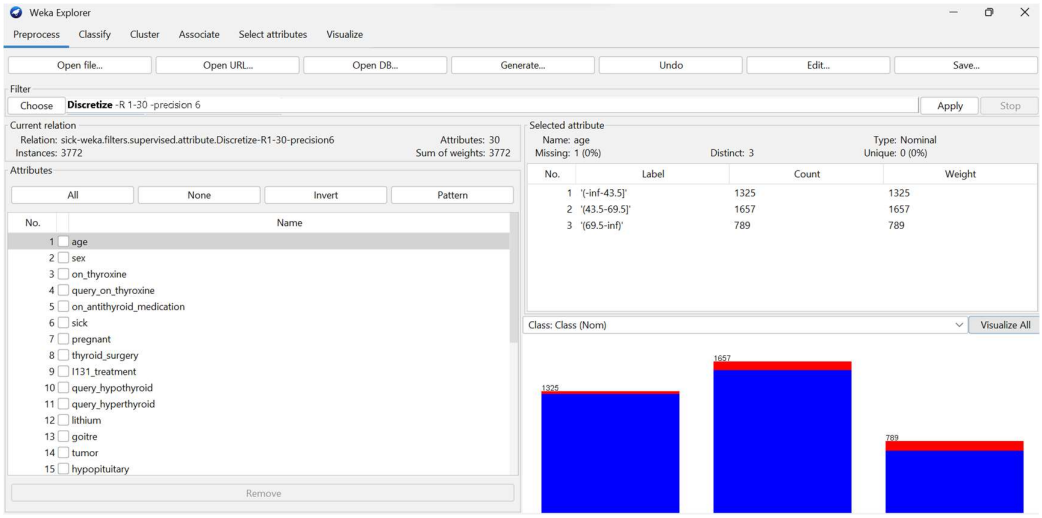
Weka Explorer
Preprocess Classify Cluster Associate Select attributes Visualize
Open file... Open URL... Open DB... Generate... Undo Edit... Save...
Filter
Choose Discretize -R 1-30 -precision 6
Apply Stop
Current relation
Relation: sick-weka.filters.supervised.attribute.Discretize-R1-30-precision6
Instances: 3772
Attributes: 30
Sum of weights: 3772
Attributes
All None Invert Pattern
No. Name
1 age
2 sex
3 on_thyroxine
4 query_on_thyroxine
5 on_antithyroid_medication
6 sick
7 pregnant
8 thyroid_surgery
9 I131_treatment
10 query_hypothyroid
11 query_hyperthyroid
12 lithium
13 goitre
14 tumor
15 hypopituitary
Remove
Selected attribute
Name: age
Missing: 1 (0%)
Distinct: 3
Type: Nominal
Unique: 0 (0%)
No. Label Count Weight
1 '(-inf-43.5]' 1325 1325
2 '(43.5-69.5]' 1657 1657
3 '(69.5-inf)' 789 789
Class: Class (Nom)
Visualize All
1325
1657
789

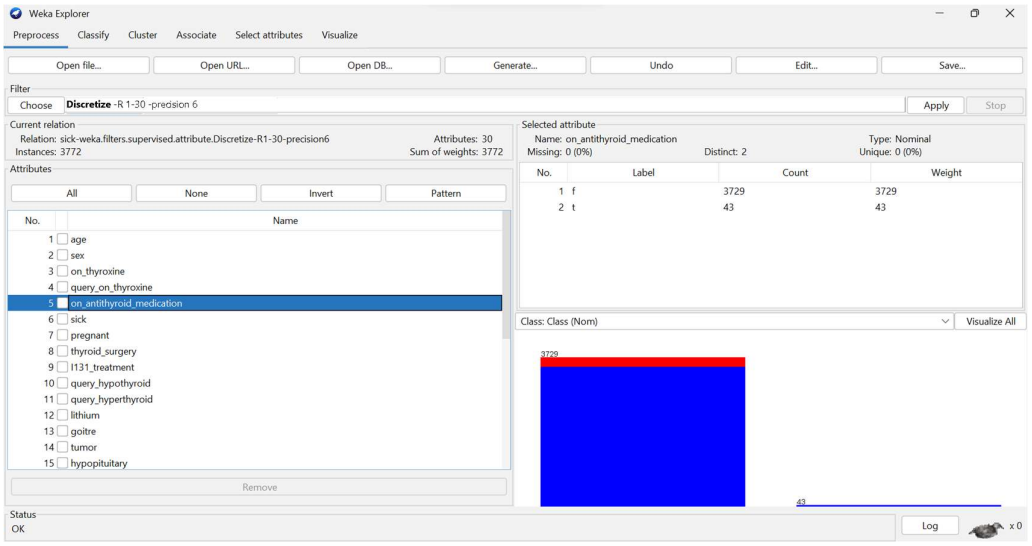
Weka Explorer
Preprocess Classify Cluster Associate Select attributes Visualize
Open file... Open URL... Open DB... Generate... Undo Edit... Save...
Filter
Choose Discretize -R 1-30 -precision 6
Apply Stop
Current relation
Relation: sick-weka.filters.supervised.attribute.Discretize-R1-30-precision6
Instances: 3772
Attributes: 30
Sum of weights: 3772
Attributes
All None Invert Pattern
No. Name
1 age
2 sex
3 on_thyroxine
4 query_on_thyroxine
5 on_antithyroid_medication
6 sick
7 pregnant
8 thyroid_surgery
9 I131_treatment
10 query_hypothyroid
11 query_hyperthyroid
12 lithium
13 goitre
14 tumor
15 hypopituitary
Remove
Selected attribute
Name: on_antithyroid_medication
Missing: 0 (0%)
Distinct: 2
Type: Nominal
Unique: 0 (0%)
No. Label Count Weight
1 f 3729 3729
2 t 43 43
Class: Class (Nom)
Visualize All
3729
43
Status
OK
Log
x 0

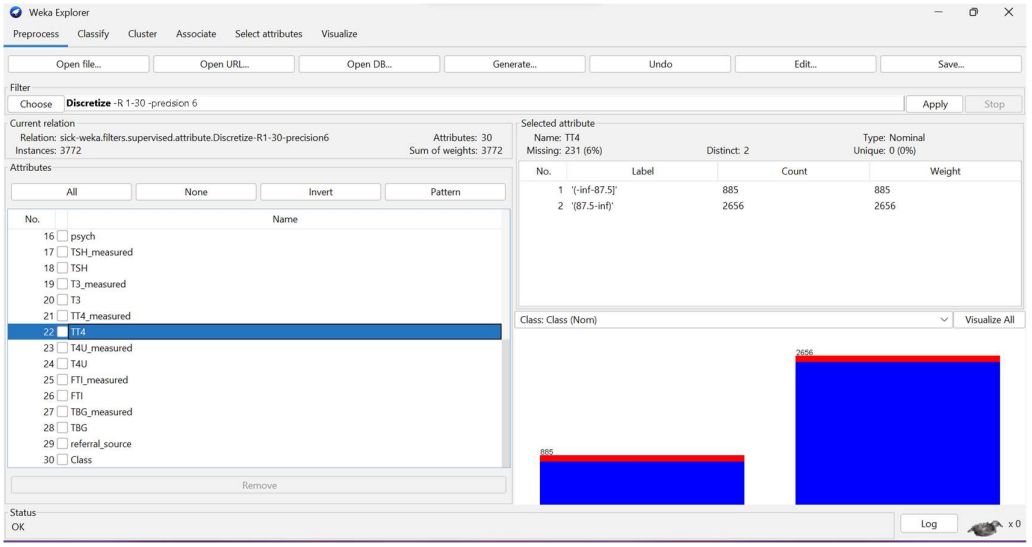
Weka Explorer
Preprocess Classify Cluster Associate Select attributes Visualize
Open file... Open URL... Open DB... Generate... Undo Edit... Save...
Filter
Choose Discretize -R 1-30 -precision 6
Apply Stop
Current relation
Relation: sick-weka.filters.supervised.attribute.Discretize-R1-30-precision6
Instances: 3772
Attributes: 30
Sum of weights: 3772
Attributes
All None Invert Pattern
No. Name
16 psych
17 TSH_measured
18 TSH
19 T3_measured
20 T3
21 TT4_measured
22 TT4
23 T4U_measured
24 T4U
25 FTI_measured
26 FTI
27 TBG_measured
28 TBG
29 referral_source
30 Class
Remove
Selected attribute
Name: TT4
Missing: 231 (6%)
Distinct: 2
Type: Nominal
Unique: 0 (0%)
No. Label Count Weight
1 '(-inf-87.5]' 885 885
2 '(87.5-inf)' 2656 2656
Class: Class (Nom)
Visualize All
885
2656
Status
OK
Log
x 0

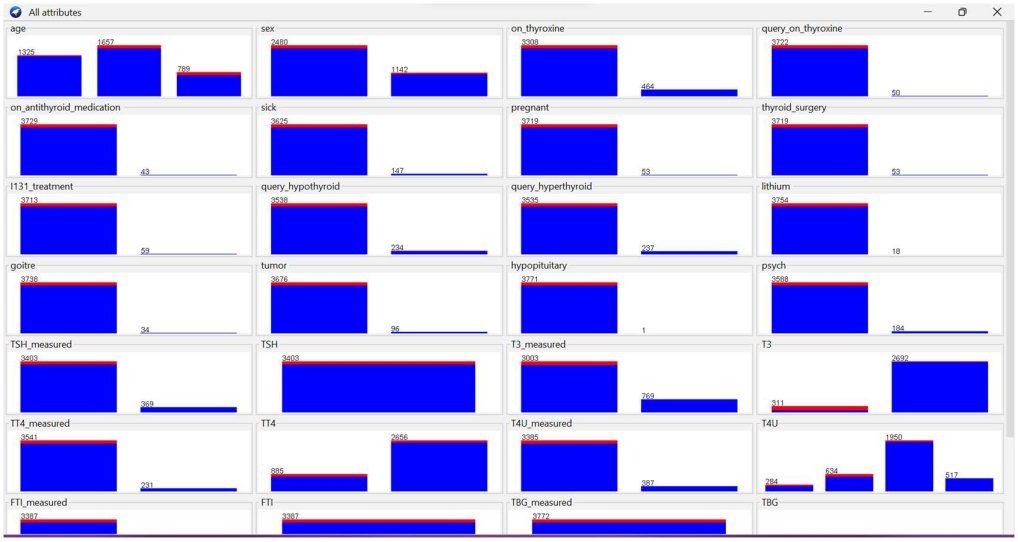

3. What is the effect of this filter on the attributes?

Ans: The discrete class intervals are formed and the frequency is calculated.

Selected attribute

Name: TT4 Type: Nominal
Missing: 231 (6%) Distinct: 2 Unique: 0 (0%)

No.	Label	Count	Weight
1	'(-inf-87.5]'	885	885
2	'(87.5-inf)'	2656	2656

Selected attribute

Name: age Type: Nominal
Missing: 1 (0%) Distinct: 3 Unique: 0 (0%)

No.	Label	Count	Weight
1	'(-inf-43.5]'	1325	1325
2	'(43.5-69.5]'	1657	1657
3	'(69.5-inf)'	789	789

4. How many distinct ranges have been created for each attribute?

Ans:

age	3
sex	2
on_thyroxine	2
query_on_thyroxine	2
on_antithyroid_medication	2
sick	2
pregnant	2
thyroid_surgery	2
I131_treatment	2
query_hypothyroid	2

query_hyperthyroid	2
Lithium	2
Goitre	2
tumor	2
hypopituitary	2
psych	2
TSH_measured	2
TSH	1
T3_measured	2
T3	1
TT4_measured	2
TT4	1
T4U_measured	2
T4U	4
FTI_measured	2
FTI	1
TBG_measured	1
TBG	1
referral_source	5
Class	2

5. Undo the filter applied in the previous step.

Ans:

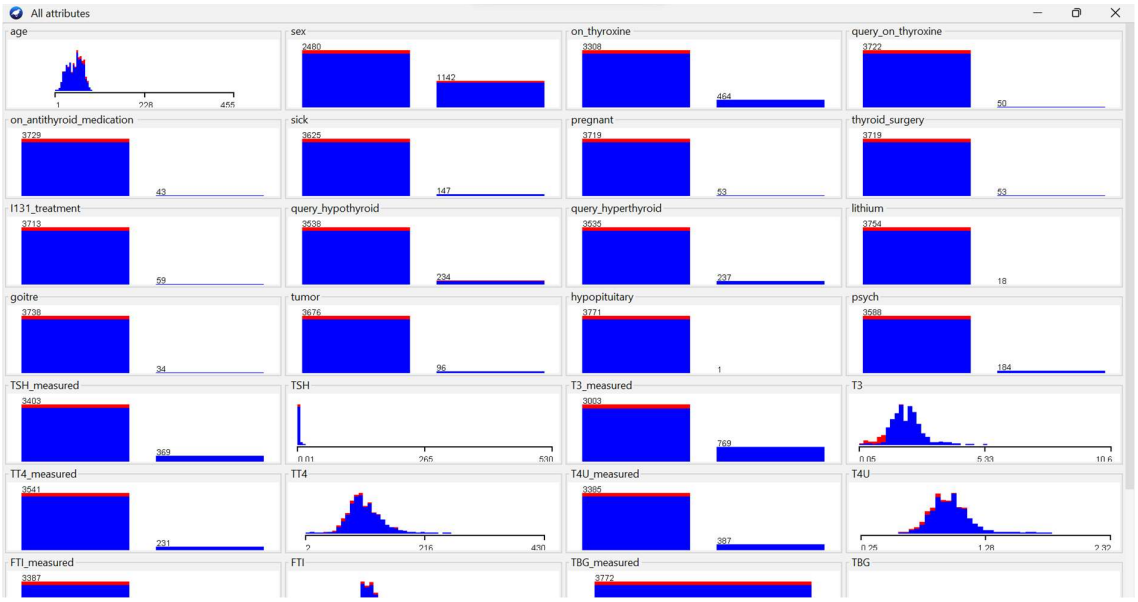

6. Apply the unsupervised discretization filter. Do this twice:

    1. In this step, set 'bins'=5

    Ans:

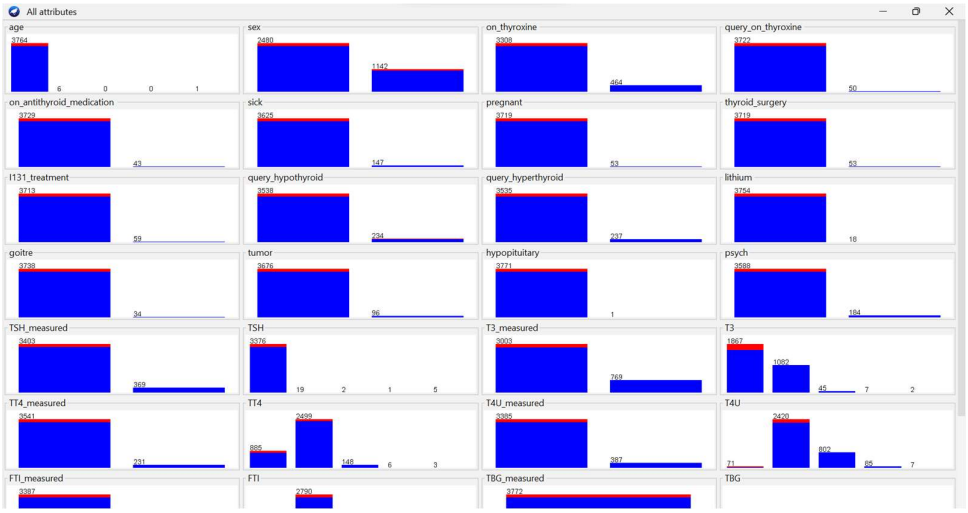

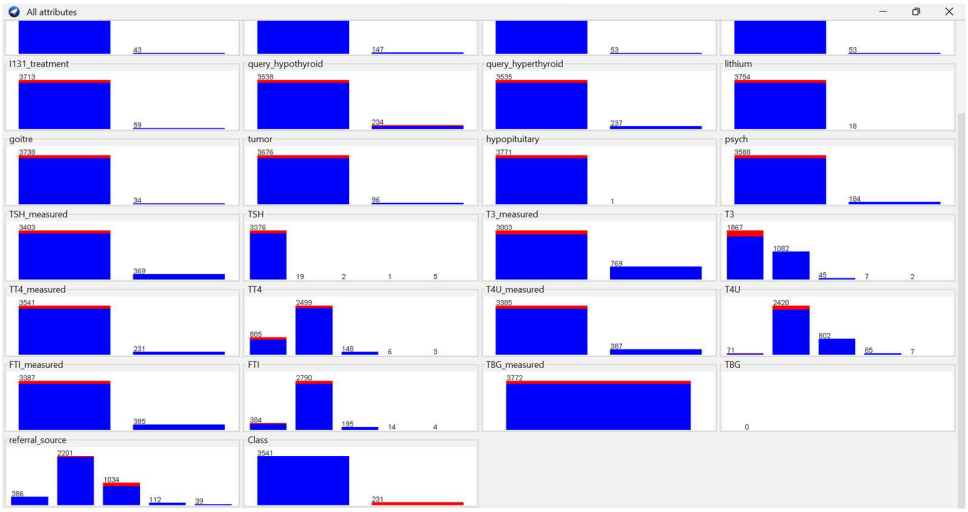

    2. In this step, set 'bins'=10

    Ans:

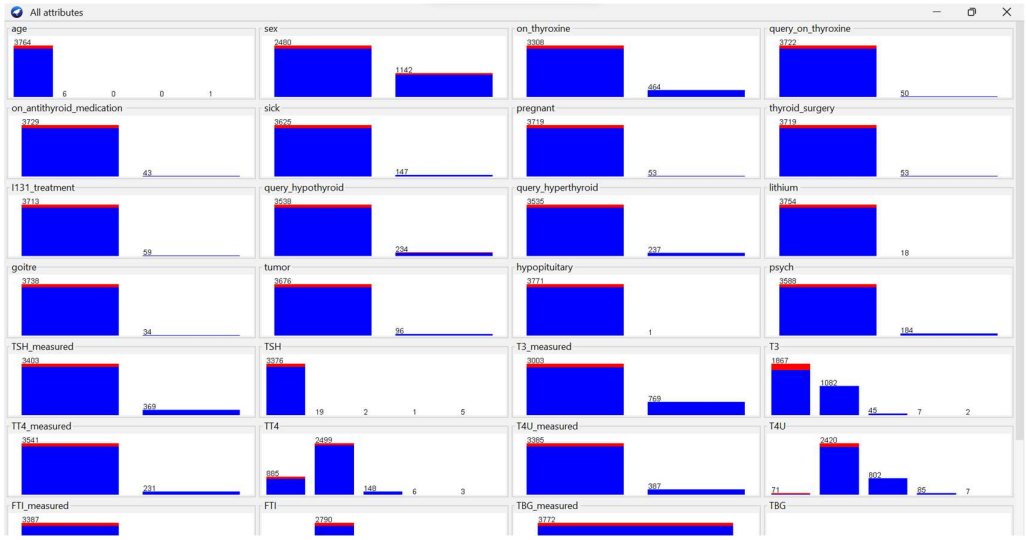

3. What is the effect of the unsupervised filter on the dataset?

Ans: Unsupervised filter work without taking any class distributions into account. The unsupervised *discretize* filter only considers the attribute being discretized. While it can 'optimize' the number of bins, it does so only with respect to self-encoding.

7. Run the Naive Bayes classifier after apply the following filters

1. Unsupervised discretized with 'bins'=5

Ans:

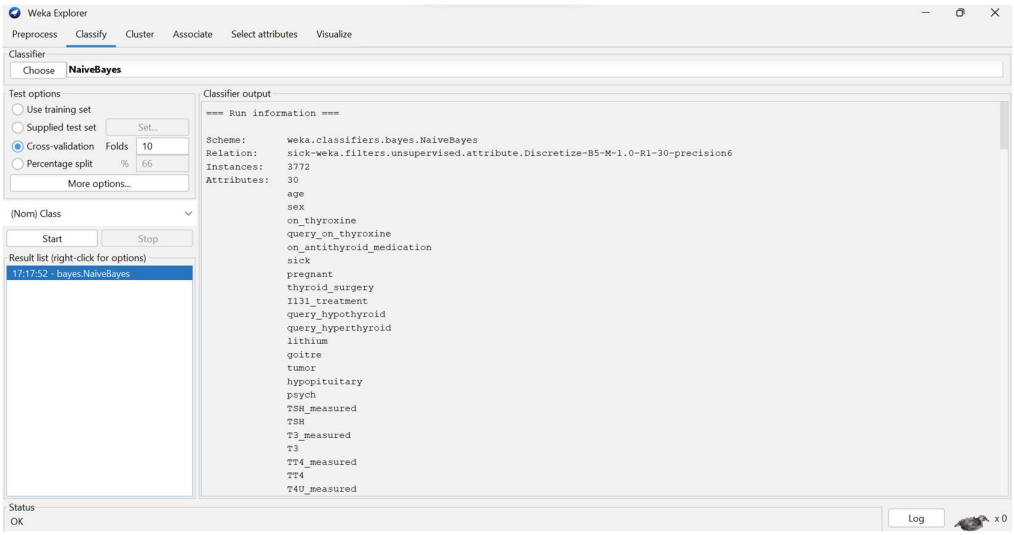

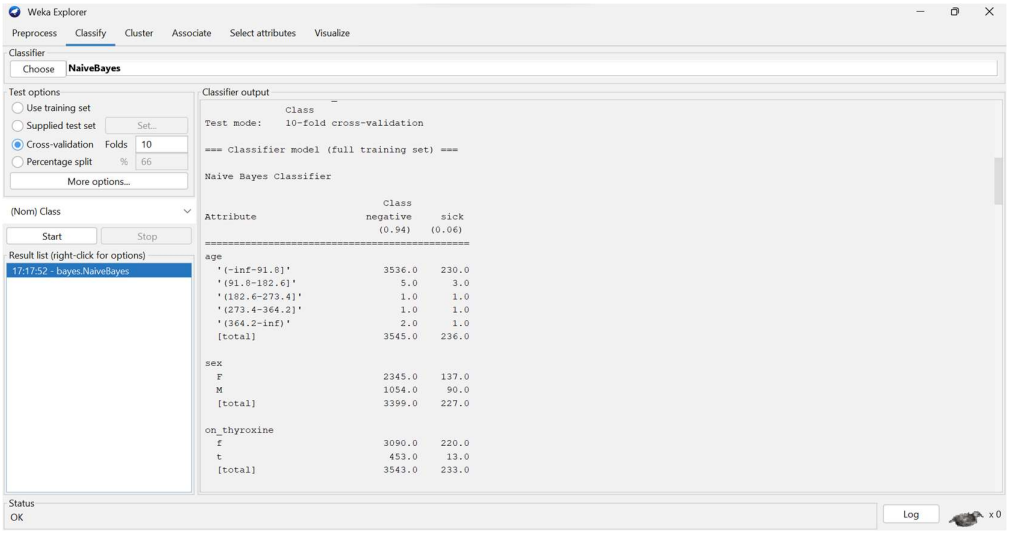

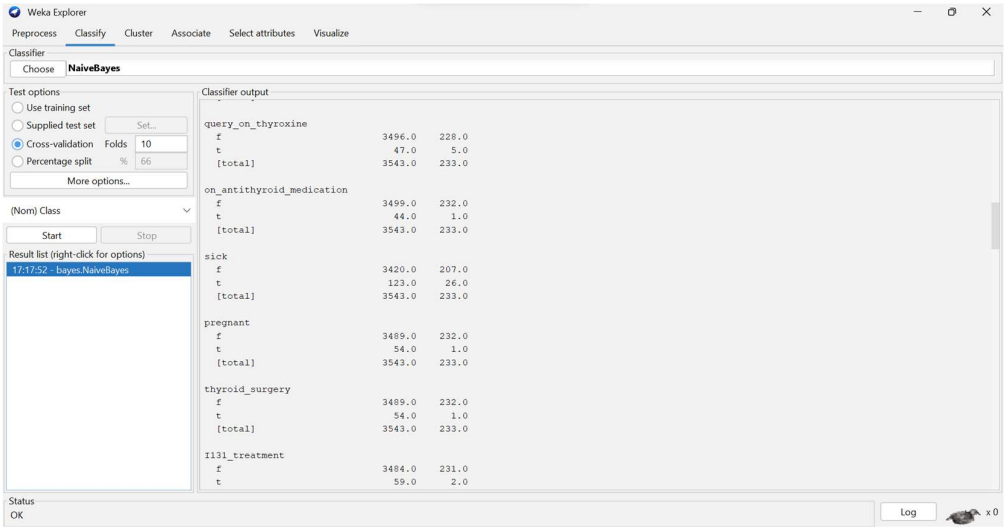
Weka Explorer
Preprocess
Classify
Cluster
Associate
Select attributes
Visualize
Classifier
Choose
NaiveBayes
Test options
Use training set
Supplied test set
Set...
Cross-validation
Folds
10
Percentage split
%
66
More options...
(Nom) Class
Start
Stop
Result list (right-click for options)
17:17:52 - bayes.NaiveBayes
Classifier output
query_on_thyroxine
f 3496.0 228.0
t 47.0 5.0
[total] 3543.0 233.0
on_antithyroid_medication
f 3499.0 232.0
t 44.0 1.0
[total] 3543.0 233.0
sick
f 3420.0 207.0
t 123.0 26.0
[total] 3543.0 233.0
pregnant
f 3489.0 232.0
t 54.0 1.0
[total] 3543.0 233.0
thyroid_surgery
f 3489.0 232.0
t 54.0 1.0
[total] 3543.0 233.0
I131_treatment
f 3484.0 231.0
t 59.0 2.0
Status
OK
Log
x 0

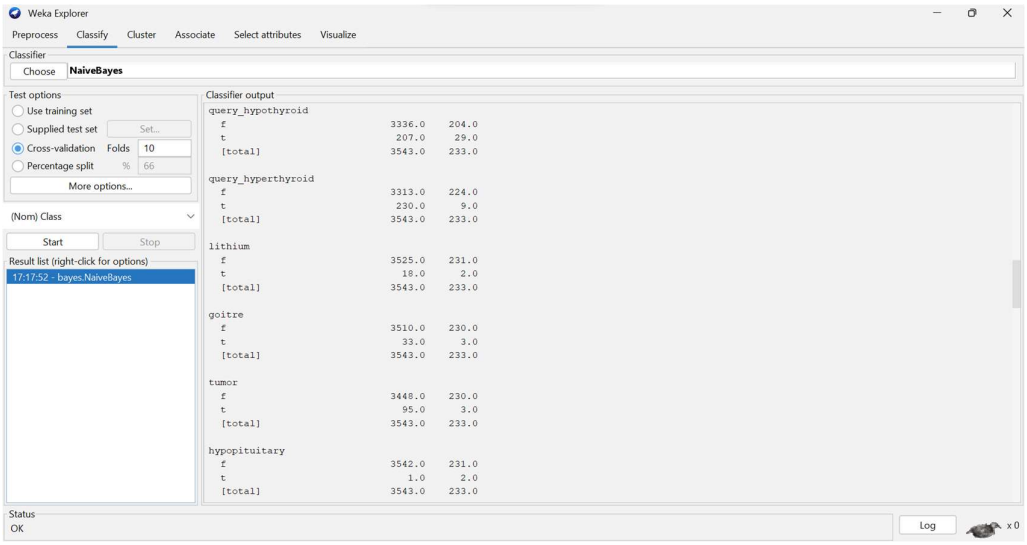
Weka Explorer
Preprocess
Classify
Cluster
Associate
Select attributes
Visualize
Classifier
Choose
NaiveBayes
Test options
Use training set
Supplied test set
Set...
Cross-validation
Folds
10
Percentage split
%
66
More options...
(Nom) Class
Start
Stop
Result list (right-click for options)
17:17:52 - bayes.NaiveBayes
Classifier output
query_hypothyroid
f 3336.0 204.0
t 207.0 29.0
[total] 3543.0 233.0
query_hyperthyroid
f 3313.0 224.0
t 230.0 9.0
[total] 3543.0 233.0
lithium
f 3525.0 231.0
t 18.0 2.0
[total] 3543.0 233.0
goitre
f 3510.0 230.0
t 33.0 3.0
[total] 3543.0 233.0
tumor
f 3448.0 230.0
t 95.0 3.0
[total] 3543.0 233.0
hypopituitary
f 3542.0 231.0
t 1.0 2.0
[total] 3543.0 233.0
Status
OK
Log
x 0

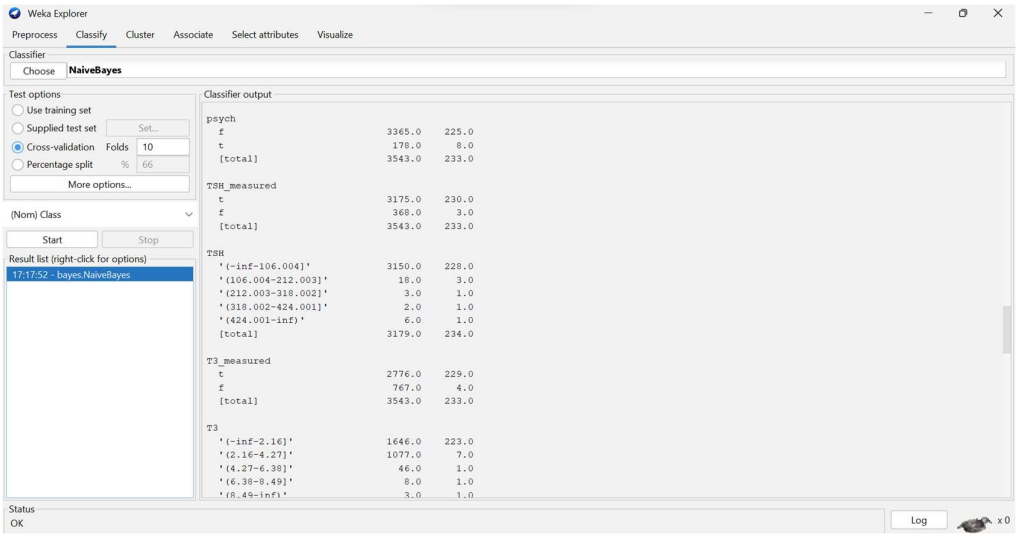
Weka Explorer
Preprocess
Classify
Cluster
Associate
Select attributes
Visualize
Classifier
Choose
NaiveBayes
Test options
Use training set
Supplied test set
Set...
Cross-validation
Folds
10
Percentage split
%
66
More options...
(Nom) Class
Start
Stop
Result list (right-click for options)
17:17:52 - bayes.NaiveBayes
Classifier output
psych
f 3365.0 225.0
t 178.0 8.0
[total] 3543.0 233.0
TSH_measured
t 3175.0 230.0
f 368.0 3.0
[total] 3543.0 233.0
TSH
'(-inf-106.004]' 3150.0 228.0
'(106.004-212.003]' 18.0 3.0
'(212.003-318.002]' 3.0 1.0
'(318.002-424.001]' 2.0 1.0
'(424.001-inf)' 6.0 1.0
[total] 3179.0 234.0
T3_measured
t 2776.0 229.0
f 767.0 4.0
[total] 3543.0 233.0
T3
'(-inf-2.16]' 1646.0 223.0
'(2.16-4.27]' 1077.0 7.0
'(4.27-6.38]' 46.0 1.0
'(6.38-8.49]' 8.0 1.0
Status
OK
Log
x 0

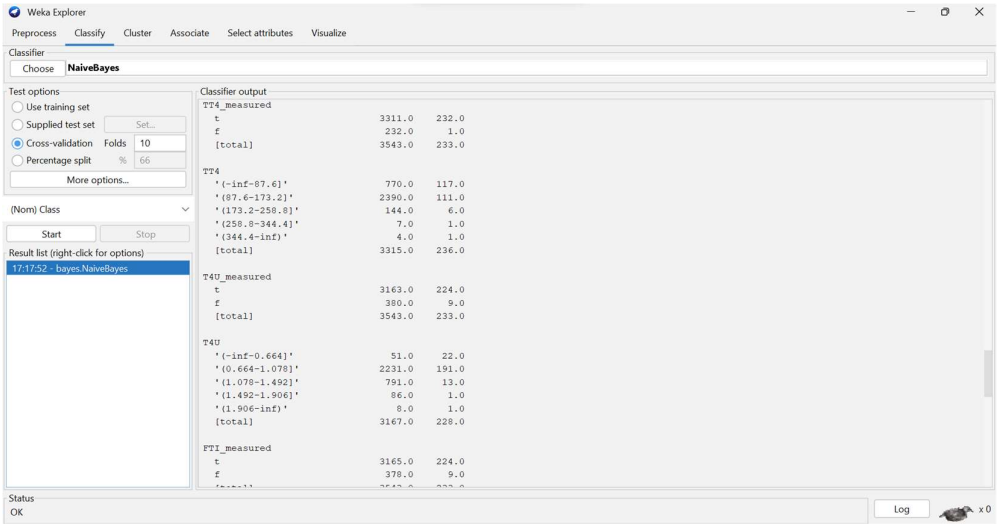

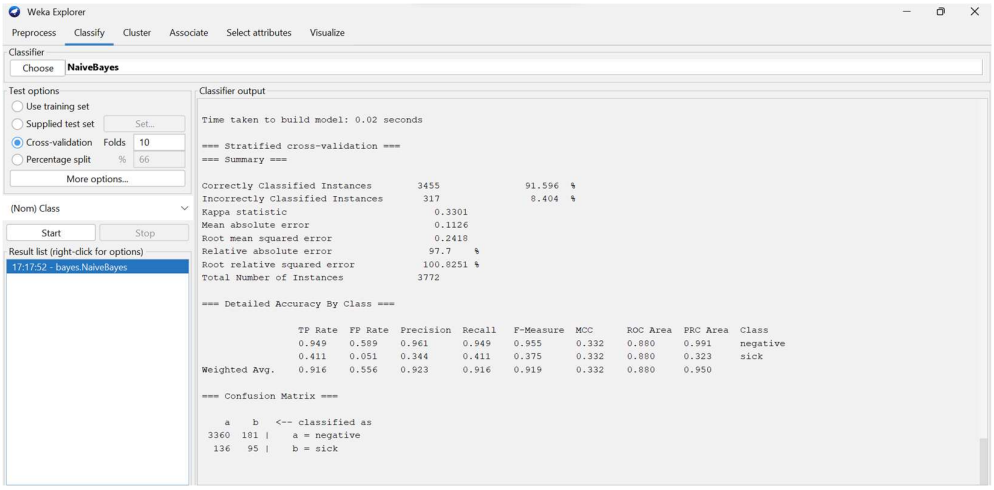

2. Unsupervised discretized with 'bins'=10

Ans:

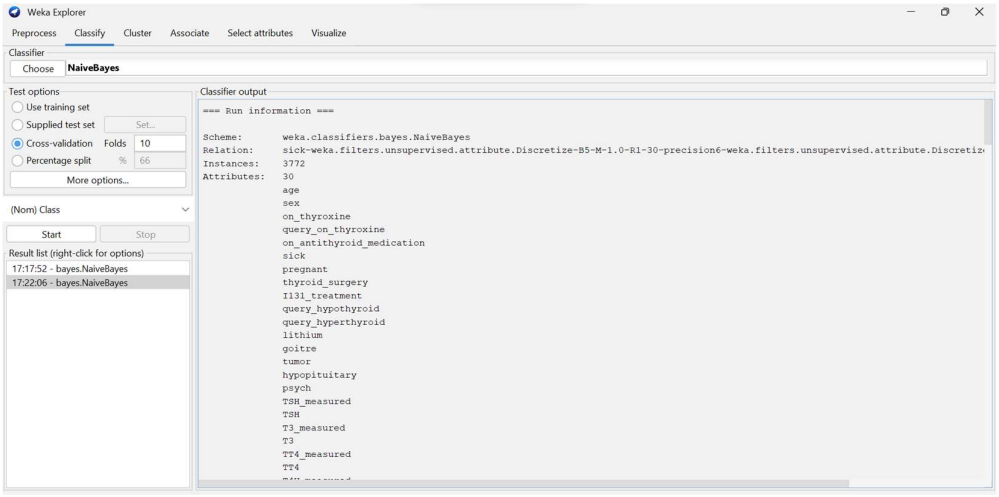

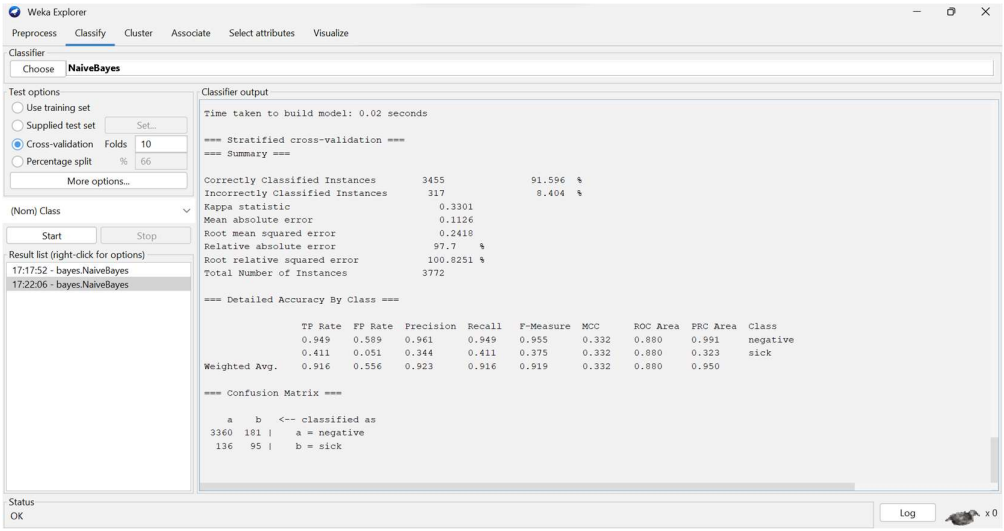

3. Unsupervised discretized with 'bins"=20.

Ans:

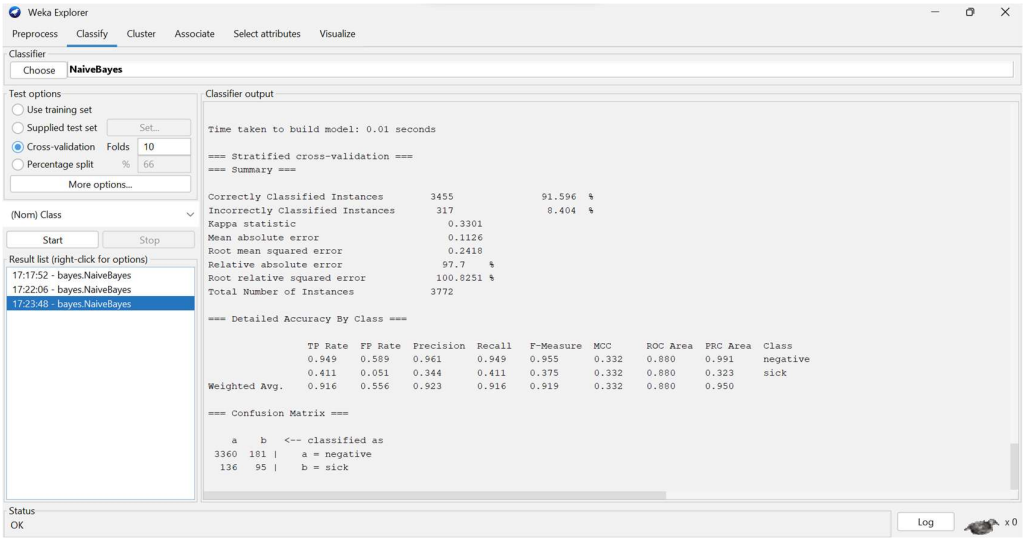

8. Compare the accuracy of the following cases

1. Naive Bayes without discretization filters

Ans:

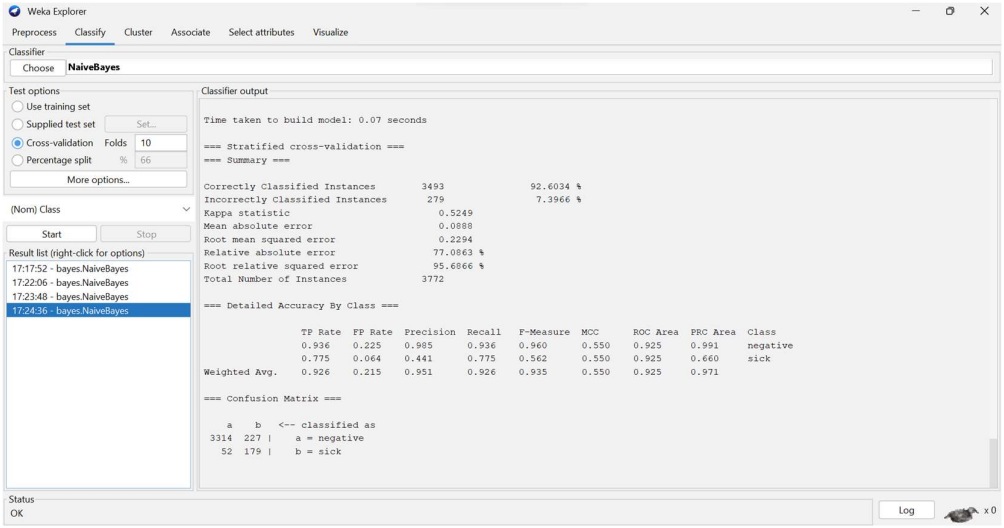

2. Naive Bayes with a supervised discretization filter

Ans:

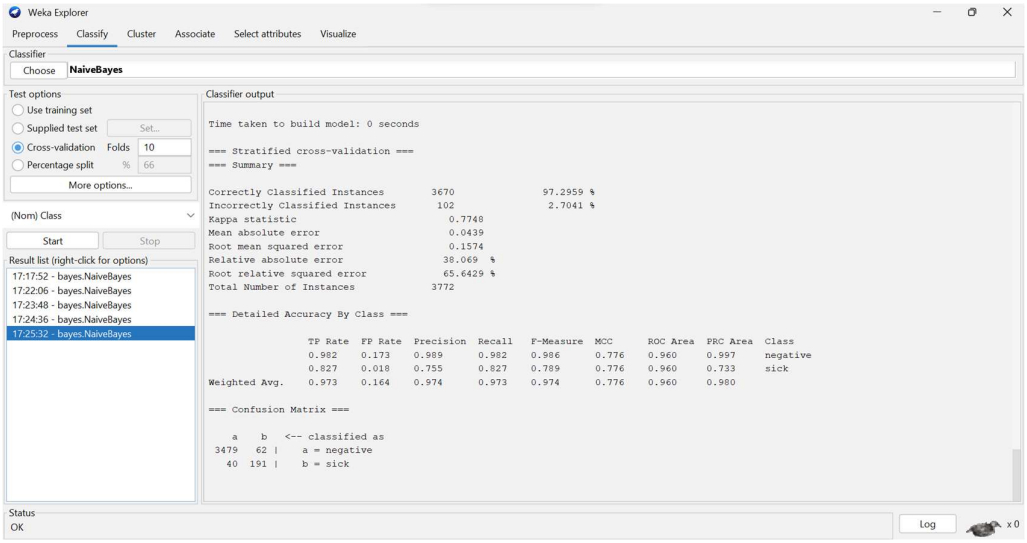

3. Naive Bayes with an unsupervised discretization filter with different values for the 'bins attributes.

Ans:

Bins = 5

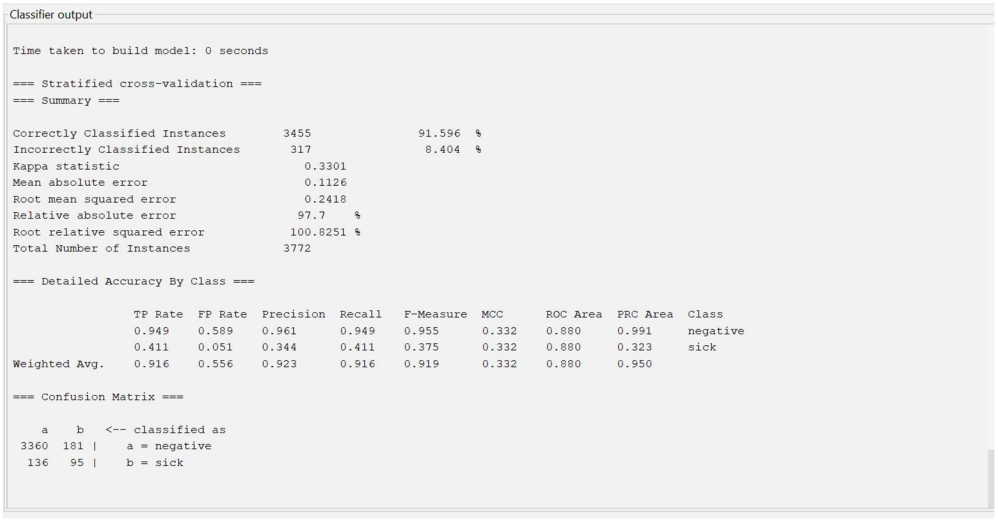

Bins = 10

Classifier output

```
Time taken to build model: 0 seconds

=== Stratified cross-validation ===
=== Summary ===

Correctly Classified Instances        3455               91.596  %
Incorrectly Classified Instances       317                8.404  %
Kappa statistic                          0.3301
Mean absolute error                      0.1126
Root mean squared error                  0.2418
Relative absolute error                 97.7    %
Root relative squared error            100.8251 %
Total Number of Instances             3772

=== Detailed Accuracy By Class ===

                 TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Class
                 0.949    0.589    0.961      0.949    0.955      0.332    0.880     0.991     negative
                 0.411    0.051    0.344      0.411    0.375      0.332    0.880     0.323     sick
Weighted Avg.    0.916    0.556    0.923      0.916    0.919      0.332    0.880     0.950

=== Confusion Matrix ===

    a    b   <-- classified as
 3360  181 |    a = negative
  136   95 |    b = sick
```

Bins = 20

Classifier output

```
Time taken to build model: 0 seconds

=== Stratified cross-validation ===
=== Summary ===

Correctly Classified Instances        3455               91.596  %
Incorrectly Classified Instances       317                8.404  %
Kappa statistic                          0.3301
Mean absolute error                      0.1126
Root mean squared error                  0.2418
Relative absolute error                 97.7    %
Root relative squared error            100.8251 %
Total Number of Instances             3772

=== Detailed Accuracy By Class ===

                 TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Class
                 0.949    0.589    0.961      0.949    0.955      0.332    0.880     0.991     negative
                 0.411    0.051    0.344      0.411    0.375      0.332    0.880     0.323     sick
Weighted Avg.    0.916    0.556    0.923      0.916    0.919      0.332    0.880     0.950

=== Confusion Matrix ===

    a    b   <-- classified as
 3360  181 |    a = negative
  136   95 |    b = sick
```